

# IFC Inside: A General Approach to Retrofitting Languages with Dynamic Information Flow Control

Stefan Heule, Deian Stefan, Edward Z. Yang,  
John C. Mitchell, Alejandro Russo

Stanford University, Chalmers University

# What is IFC?


- Not all information in a program is equal
  - Some might be more sensitive
- Information flow control ...
  - ... *tracks* where information flows
  - ... allows *policies to restrict* flows of information

# Example: security on the web

- Modern web pages involve many components
  - Javascript of the website
  - Javascript of some (untrusted) library
  - Advertising code
  - Browser addons
- Web content can be very sensitive
  - Online banking, passwords, personal information, etc.

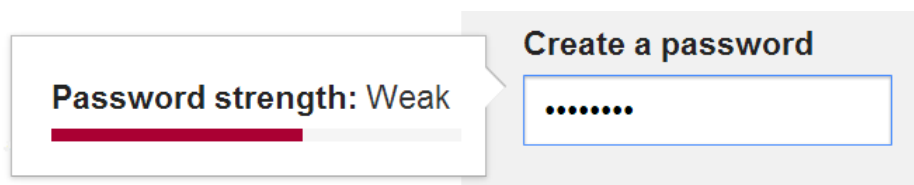
# Current situation

- Code written by many different parties
  - Potentially mutually distrusting parties
  - Computing over sensitive data
- IFC to the rescue
  - Label sensitive data as such
  - Prevent flow of sensitive data to undesired places (like arbitrary web servers)



# Concrete Example

- Password strength checking



Password strength: Weak

Create a password

.....


- Website uses `check_strength(pw)` from a (not fully trusted) library
  - The library could send the password to `bad.com`
- With IFC, we can precisely decide what the library can do with the password

# Why isn't everyone using IFC?

- We are interested in a *dynamic* IFC system
- A key challenge is performance
  - Tracking information at a fine-grained level is expensive

# Coarse-grained IFC

- The program is split into computational units (tasks)
  - All data within one task has a single label
- Different computational units can communicate



# Advantages of coarse-grained IFC

1. Efficiency
  - Checks only at isolation boundaries
2. Minimal changes to language
  - If added to a language retroactively
3. Reuse existing programs
4. Simple
  - to understand and reason about



# Goals

- Formally define a core calculus for a coarse-grained dynamic IFC system
- Combine IFC language with *any* programming language
- Prove security guarantees of IFC (known as non-interference)

# Approach Overview

- Given a **target language** (any programming language)
- Define an **IFC language**
  - Minimal calculus, only IFC features
- Combine **target** and **IFC** language
  - Allow **target** language to call into **IFC**, and vice-versa
- Careful definition of the IFC language allows the overall system to provide isolation, regardless of what the target language does

# IFC language

- Tag data with security labels
  - Labels form a lattice, and determine how data can flow inside an application
- Example lattice
  - Two labels  $H$  (high) and  $L$  (low)
  - Flow from  $H$  to  $L$  is not allowed

$H$   
↑  
 $L$

# IFC language: labels


- Get and set the current label
  - **setLabel**, **getLabel**



- Setting the label is only allowed to *raise* the label
- Can also compute on labels
  - $\sqsubseteq, \sqcap, \sqcup$


# IFC language: sandboxing

- Isolate an expression as a new task
  - **sandbox  $e$**



# Inter-task communication


- Tasks can send and receive messages
- Send message  $\mathbf{v}$  to task  $\mathbf{i}$ , protected by label  $\mathbf{l'}$ 
  - **send  $\mathbf{i} \mathbf{l' v}$**
  - Can only send messages at or above current label



# Inter-task communication

- Receiving either binds a message **v** and sender **i** in  **$e_1$** , or execution continues in  **$e_2$**  (if there is no message)
  - Messages that are above the current level are never received

**recv  $i, v$  in  $e_1$  else  $e_2$**



# Formal treatment



# What is a programming language?

- Need a formal definition of a language
  - Global store  $\Sigma$
  - Evaluation context  $E$
  - Expression syntax  $e$ , some expressions are values  $v$
  - Reduction relation  $\rightarrow$
- This is the **target language**

# Example: Mini-ECMAScript

$v ::= \lambda x.e \mid \text{true} \mid \text{false} \mid a$   
 $e ::= v \mid x \mid e e \mid \text{if } e \text{ then } e \text{ else } e$   
 $\quad \mid \text{ref } e \mid !e \mid e := e \mid \text{fix } e$   
 $E ::= [\cdot]_T \mid E e \mid v E \mid \text{if } E \text{ then } e \text{ else } e$   
 $\quad \mid \text{ref } E \mid !E \mid E := e \mid v := E \mid \text{fix } E$

T-APP

$$\frac{}{\mathcal{E}_\Sigma [(\lambda x.e) v] \rightarrow \mathcal{E}_\Sigma [\{v / x\} e]}$$

T-IFTRUE

$$\frac{}{\mathcal{E}_\Sigma [\text{if } e_2] \rightarrow \mathcal{E}_\Sigma [e_1]}$$

T-IFFALSE

$$\frac{}{\mathcal{E}_\Sigma [\text{if false then } e_1 \text{ else } e_2] \rightarrow \mathcal{E}_\Sigma [e_2]}$$

T-REF

$$\frac{\text{fresh}(a)}{\mathcal{E}_\Sigma [\text{ref } v] \rightarrow \mathcal{E}_{\Sigma[a \mapsto v]} [a]}$$

T-DEREF

$$\frac{(a, v) \in \Sigma}{\mathcal{E}_\Sigma [!a] \rightarrow \mathcal{E}_\Sigma [v]}$$

T-ASS

$$\frac{}{\mathcal{E}_\Sigma [a := v] \rightarrow \mathcal{E}_{\Sigma[a \mapsto v]} [v]}$$

T-FIX

$$\frac{}{\mathcal{E}_\Sigma [\text{fix } (\lambda x.e)] \rightarrow \mathcal{E}_\Sigma [\{\text{fix } (\lambda x.e) / x\} e]}$$

# Notation

- Rules are standard, except we use  $\mathcal{E}_{\Sigma}$  instead of normal context  $\mathbf{E}$

$$\frac{\text{T-IFFALSE}}{\mathcal{E}_{\Sigma} [\text{if false then } \mathbf{e}_1 \text{ else } \mathbf{e}_2] \rightarrow \mathcal{E}_{\Sigma} [\mathbf{e}_2]}$$

- Obtain normal semantics with


$$\mathcal{E}_{\Sigma} [\mathbf{e}] \triangleq \Sigma, \mathbf{E} [\mathbf{e}]$$

- Later, we re-interpret what  $\mathcal{E}$  stands for

# IFC language

- Also defined in terms of an special  $\mathcal{E}$

$$\frac{\text{I-SETLABEL} \quad l \sqsubseteq l'}{\mathcal{E}_{\Sigma}^{i,l} [\text{setLabel } l'] \rightarrow \mathcal{E}_{\Sigma}^{i,l'} [\langle \rangle]}$$



# Embedding [Matthews and Findler, POPL'07]

- Extend IFC and target language syntax

$$\begin{aligned} v &::= \dots \mid \text{IT}[\mathbf{v}] \\ e &::= \dots \mid \text{IT}[\mathbf{e}] \\ E &::= \dots \mid \text{IT}[\mathbf{E}] \end{aligned}$$

$$\begin{aligned} \mathbf{v} &::= \dots \mid \text{TI}[v] \\ \mathbf{e} &::= \dots \mid \text{TI}[e] \\ \mathbf{E} &::= \dots \mid \text{TI}[E] \end{aligned}$$

- Re-interpret context and reduction relation

$$\begin{aligned} \mathcal{E}_{\Sigma}[\mathbf{e}] &\triangleq \Sigma; \langle \Sigma, E[\mathbf{e}]_{\mathbf{T}} \rangle_l^i, \dots \\ \mathcal{E}_{\Sigma}^{i,l}[e] &\triangleq \Sigma; \langle \Sigma, E[e]_I \rangle_l^i, \dots \end{aligned}$$

# Sandboxing


I-SANDBOX

$$\begin{array}{c}
 \Sigma' = \Sigma \left[ i' \mapsto \epsilon \right] \quad \Sigma' = \kappa (\Sigma) \\
 t_1 = \langle \Sigma, E[i'] \rangle_l^i \quad t_{\text{new}} = \langle \Sigma', e \rangle_l^{i'} \quad \text{fresh}(i') \\
 \hline
 \Sigma; \langle \Sigma, E[\text{sandbox } e]_I \rangle_l^i, \dots \xrightarrow{\alpha} \Sigma'; \alpha_{\text{sandbox}}(t_1, \dots, t_{\text{new}})
 \end{array}$$

- **sandbox e** does
  - Create new task for e
  - Schedule e according to scheduling policy

# Security Guarantees

- Non-interference:
  - Intuitively: An attacker that can only see values up to level  $l$  should not see a difference in behavior if values at level  $l' > l$  are changed



# Erasure function

- Formally, we need an erasure function  $\varepsilon_l$ 
  - Erases all data above  $l$  to ■
  - Program  $c_1$  and  $c_2$  are  $l$ -equivalent,  $c_1 \approx_l c_2$ , iff  $\varepsilon_l(c_1) = \varepsilon_l(c_2)$
- For our system,  $\varepsilon_l$  erases the following:
  - Any tasks with current label above  $l$
  - Any messages with label above  $l$



# Termination sensitive non-interference (TSNI)

For all programs  $c_1, c_2, c'_1$  and labels  $l$ , such that

$$c_1 \approx_l c_2 \quad \text{and} \quad c_1 \hookrightarrow^* c'_1$$

then there exists  $c'_2$  such that

$$c'_1 \approx_l c'_2 \quad \text{and} \quad c_2 \hookrightarrow^* c'_2$$

**Theorem:** Any target language combined with our IFC language with round robin scheduling satisfies TSNI.

# Real world examples

- Is this actually practical?
- One challenge are external effects
  - File system, internet connection, etc.
- Possible solutions
  - Make external effects inaccessible
  - Internalize them into the IFC language
    - Labeled file system

# SWAPI

- Implementation of coarse-grained dynamic IFC system for Javascript
- Websites have access to XHR constructor
  - XHR requests need to be modeled at the IFC language level
- SWAPI chooses origins (`example.com`) as labels

# Password strength checker in SWAPI

- Even if `check_strength(pw)` is completely untrusted, it cannot send the password to `bad.com`
- Execute `check_strength` in a sandboxed task

# Conclusions

- Coarse-grained IFC is great
  - Allows for language-independent IFC system
  - Efficient, yet flexible
- Combining operational semantics of two languages as key mechanism to formalize our system

Thank you.

**Questions?**

